



AdPilot OS

A read-only operating system for paid ads. It audits your Meta & TikTok ad accounts, scores a 0–100 Campaign Health Score, and proposes safe, numbers-first fixes — it never edits, pauses, creates, or spends on a live ad without your explicit approval. **It proposes; the human approves.**

v6.0.0 · Australian English & AUD · 713 automated tests green · Next.js 16 · React 19 · Supabase (RLS) · Stripe

1 Campaign Health Score

0–100, computed from 13 weighted factors. Factors with no data redistribute their weight to the rest, so the score stays fair. Four bands set the next action:

GREEN
≥ 80

YELLOW
≥ 60

ORANGE
≥ 40

RED
< 40

Green = scale-eligible (if tracking is clean) · Yellow = fix weak factors first · Orange = act this week · Red = stop & fix before spending more.

Tracking quality		15
CPA vs break-even		15
Spend efficiency		12
Conversion rate		10
CTR		8
Lead quality		8
Creative freshness		8
CPC		7
Naming quality		5
Offer strength		5
Landing-page alignment		4
Budget pacing		2
Data confidence		1

2 Per-ad decisions — seven verdicts

Every ad gets exactly one verdict. **Scale** and **kill** only fire when a **Wilson significance test** confirms the result — never on a noisy sample. Each verdict carries a safe, reversible proposal.

SCALECPA ≤ break-even, health ≥ 70, tracking clean, win significant → +≤20% budget & duplicate

KILLCPA > 1.5× break-even, loss significant → pause & reallocate

REDUCECPA moderately high → cut budget, test a fresh angle

REFRESHFrequency ≥ 4 with falling or low CTR → 3–5 fresh variants

FIX-TRACKINGSpend with zero/implausible results → audit pixel / events / CAPI first

KEEPWithin range → monitor, no change

INSUFFICIENT< 50 clicks & < 15 conversions → let spend accumulate

Lead-gen uses a break-even-CPL split (optional lead_close_rate); never auto-kills on leads.

The read-only invariant — proposes, never acts on its own

AdPilot is **read-only by default**: it never edits, pauses, creates, or spends on a live ad. Acting on a live ad requires an **explicit typed “YES”**, is **Expert-tier only**, and is double-gated by the environment flag `ADS_WRITE_ENABLED` — which ships **OFF**. Every engine decision is marked `safe`; all connections use read-only API scopes; numbers-first, anti-hype — no guarantees and no financial, legal, medical or tax advice.

3 Twelve AI specialists

A **Command** router sends each question to the right expert. All are grounded in product knowledge, share hardened guardrails, and are read-only & numbers-first (no invented figures). Light creative work is cost-routed to a smaller model; the system prompt is cached.

 **Command** — routes & prioritises

 **Dana** — data & unit economics

 **Milo** — automation (no-code)

 **Paige** — policy & safety

 **Mira** — Meta ads

 **Stella** — creative

 **Atlas** — tracking & pixels

 **Piper** — productisation

• **Travis** — TikTok ads

 **Titan** — offer & funnel

 **Riley** — reporting (prose)

 **Quinn** — QA & sanity checks

+ Inside the workspace

Command Centre (hub: score, proposals, trends, Data Sync) · **Connect** · **Dashboard / Manual** (CSV) · **Proposals** queue · **Reports** · **Creative Studio** (asset library) · **Content + Calendar** (organic publishing) · **AI Specialists** · **Policy Check** (Meta/TikTok/ACCC copy compliance) · **Audience** (follower intelligence) · **Boost & Reach** (organic) · **Ad Actions** (Expert guarded writes) · **Messenger** (Expert) · **UTM Builder** · multi-client **Portfolio** · **Settings** · **Billing** · **Notifications**

4 Statistical diagnostics NEW · V6

- **Change-point creative-fatigue onset** — pinpoints the day engagement stepped down (falling hold-rate as a leading signal) with a confidence label, before performance craters.
- **Account-level anomaly scan** — robust median/MAD z-score flags harmful spikes & drops with watch/alert severity.
- **purchase_roas reconciliation** — Meta’s attribution-window ROAS shown beside derived ROAS as an honest cross-check.
- **Live ingestion + Data Sync panel** — append-only audit of every pull (ok / partial / rate-limited), rows, API version, time.

6 Reports & delivery

- **Five report kinds** — daily, weekly, monthly briefs + first-look Meta & TikTok audits.
- **Deterministic tables** for every number; **Riley** writes the plain-English prose only.
- Optional **projected-impact** section forecasting a proposed change.
- **Scheduled PDF delivery** — the weekly digest emails the report PDF (via Resend); Drive/Gmail delivery owner-gated.

5 Automation & safety NEW · V6

- **Configurable rules engine** — absolute + rolling-baseline (z-score / WoW) rules emit alerts & proposals; a min-volume gate kills false alarms. Never executes.
- **“What to test next” recommender** — ICE-ranked, propose-only test backlog: one-variable hypothesis, primary metric, budget, 7–14d read-out, paused-duplicate setup.
- **AI usage & cost telemetry** — per-call tokens & AUD cost, best-effort.
- **AI eval harness** — deterministic safety guards + adversarial fixtures gate the read-only invariant in CI.

7 Connect & data

- **Read-only connect** — Meta (ads_read, read_insights) & TikTok (ads.read) OAuth, a pasted System User token, or CSV import (free).
- **30-day sync window**; rate-limit backoff & per-account resilience (fail-closed).
- Tokens encrypted at rest (AES-256-GCM); every table is **RLS**-scoped to org members.
- One-click **“Run my first audit”** after connecting.

8 Two modes · four plans (AUD)

Simple mode = a 10-second snapshot; **Advanced** = full diagnostics, charts & every finding (toggle anytime). Plans gate features; annual ≈ two months free.

Free \$0 CSV import + Campaign Health Score. No card required.	Starter \$49 /mo Saved reports, content upload & publishing, threshold alerts.	Pro \$149 /mo API auto-sync, the 12 AI specialists, Creative Studio, multi-client, lead-quality loop.	Expert \$399 /mo White-label reports, Messenger automation, guarded ad changes, expert plugins.
---	--	---	---